

Tabular: Amazon (solved)

Business Context

You are a **category manager** at an Indian fashion brand selling on Amazon India. It is July 2022, end of Q1. Your data covers April–June 2022 (~120,000 orders).

You have two decisions to make before next quarter:

1. **Strategic:** Which product categories should you double down on next season — high-volume (Sets, Kurtas) or high-AOV (Western Dress)?
2. **Operational:** Your cancellation rate is ~9%. What is the single most actionable fix?

As a secondary objective, you will have built the dashboard that answers both questions.

Data Familiarization

You're working with real Amazon India Female fashion sales data from their Q2 2022 MRP which you can download [here](#). This dataset provides detailed insights into Amazon sales data, including SKU, product design, category, size, pricing, and fulfillment details. It is useful for analyzing sales performance and optimizing product profitability. Before doing any analysis, you need to understand what data you're looking at:

- What does each row represent?
- What kind of columns are available?
- Are there missing values?

First step: always **look at your data**. Use `.head()` / `.tail()` to see a sample, `.info()` to understand the structure, and `.describe()` for numeric summaries. We'll answer those questions using a few simple commands in pandas.

Typically datasets come with a *data dictionary* or a *data card* like this one:

Variable	Description	Type
category	Type of product (e.g., kurta, set, top)	String
size	Size of the product	String
date	Date of the sale	Date
status	Status of the order (e.g., Shipped, Cancelled)	String
fulfilment	Who fulfilled the order (Amazon or Merchant)	String
style	Style identifier of the product	String
sku	Stock Keeping Unit, unique product ID	String
asin	Amazon Standard Identification Number	String
courier_status	Current status of the shipping courier	String
qty	Quantity of items ordered	Integer
amount	Sale amount per item (in INR)	Float
b2b	Whether the customer is a business	Boolean
currency	Currency used for the transaction (usually INR)	String

Let's load it!

```
In [43]: import pandas as pd

### Download it:
# df = pd.read_csv('https://raw.githubusercontent.com/ciri/iese-dsfb/main/resources/tabular/AmazonSalesReport.csv')

### Use local files
df = pd.read_csv('../resources/tabular/AmazonSalesReport.csv')
```

```
/tmp/ipykernel_22712/412006718.py:7: DtypeWarning: Columns (24) have mixed types. Specify dtype option on import or set low_memory=False.
df = pd.read_csv('../resources/tabular/AmazonSalesReport.csv')
```

You may be getting a warning. When pandas loads a CSV, it tries to infer the data type of each column (int, float, object, etc.). If a column has both numbers and text, pandas gets confused. Column 23 (24th column, likely 'fulfilled-by') has mixed types - probably some strings like "Easy Ship" and some missing values (NaN) or unexpected values like False.

```
In [44]: df.describe()
df.head(3)
```

```
Out[44]:
```

	index	Order ID	Date	Status	Fulfilment	Sales Channel	ship-service-level	Style	base-SKU	SKU	...	currency	Amount	
0	0	8078784-5731545	405-30-22	Cancelled	Merchant	Amazon.in	Standard	SET389	SET389	SET389-KR-NP-S	...	INR	647.62	
1	1	9198151-1101146	171-30-22	Shipped - Delivered to Buyer	Merchant	Amazon.in	Standard	JNE3781	JNE3781	JNE3781-KR-XXXL	...	INR	406.00	B
2	2	0687676-7273146	404-30-22	Shipped	Amazon	Amazon.in	Expedited	JNE3371	JNE3371	JNE3371-KR-XL	...	INR	329.00	

3 rows × 25 columns



Pandas Quick Reference

Before diving into the cleaning and analysis, here are the three operations you'll use most — demonstrated on the data you just loaded.

Operation	What it does
<code>df[mask]</code>	Filter rows using a condition
<code>df.merge(other, on=col)</code>	Join two tables on a shared column

Filtering rows

Rows are filtered by *Boolean masks* — an array of True/False values computed from a condition. You wrap the dataframe in square brackets with the mask:

```
In [45]: f_shipped = df['Status'] == 'Shipped'
f_amazon = df['Fulfilment'] == 'Amazon'

df[f_shipped & f_amazon][['Status', 'Fulfilment', 'Category', 'Amount']].head(5)
```

```
Out[45]:
```

	Status	Fulfilment	Category	Amount
2	Shipped	Amazon	kurta	329.0
4	Shipped	Amazon	Top	574.0
5	Shipped	Amazon	Set	824.0
6	Shipped	Amazon	Set	653.0
9	Shipped	Amazon	kurta	363.0

Merging datasets

In practice, data lives in **separate systems**: order details (status, shipping) sit in the order management system; product details (category, size) sit in the product catalog. To analyze them together you need a **merge** — the pandas equivalent of Excel's VLOOKUP or a SQL JOIN.

Here we load two separate files and join them on the shared key `SKU`. A `Left` join keeps every row from the orders file — no information lost.

```
In [46]: ## Download it
# orders = pd.read_csv('https://raw.githubusercontent.com/ciri/iese-dsfb/main/resources/tabular/AmazonOrder
# products = pd.read_csv('https://raw.githubusercontent.com/ciri/iese-dsfb/main/resources/tabular/AmazonProdu

## Use local files
orders = pd.read_csv('../resources/tabular/AmazonOrders.csv')
products = pd.read_csv('../resources/tabular/AmazonProducts.csv')

orders.head(2)
products.head(2)
```

```
Out[46]:
```

	SKU	base-SKU	Style	Category	Size
0	SET389-KR-NP-S	SET389	SET389	Set	S
1	JNE3781-KR-XXXL	JNE3781	JNE3781	kurta	3XL

```
In [47]: df_merged = orders.merge(products, on='SKU', how='left')
df_merged[['Order ID', 'Status', 'Fulfilment', 'Category', 'Amount']].head(5)
```

```
Out[47]:
```

	Order ID	Status	Fulfilment	Category	Amount
0	405-8078784-5731545	Cancelled	Merchant	Set	647.62
1	171-9198151-1101146	Shipped - Delivered to Buyer	Merchant	kurta	406.00
2	404-0687676-7273146	Shipped	Amazon	kurta	329.00
3	403-9615377-8133951	Cancelled	Merchant	Western Dress	753.33
4	407-1069790-7240320	Shipped	Amazon	Top	574.00

Data Cleaning

Before we can trust or analyze this data, we need to clean it. That means:

- Standardize column names so they're consistent
- Drop unnecessary columns (like Excel export leftovers)
- Handle missing values in a simple, safe way

Drop irrelevant columns

```
In [48]: df.columns # at the end there's a weird one, likely someone added one comma too many or had a value in excel
# we also don't need this index
```

```
Out[48]: Index(['index', 'Order ID', 'Date', 'Status', 'Fulfilment', 'Sales Channel ',
               'ship-service-level', 'Style', 'base-SKU', 'SKU', 'Category', 'Size',
               'ASIN', 'Courier Status', 'Qty', 'currency', 'Amount', 'ship-city',
               'ship-state', 'ship-postal-code', 'ship-country', 'promotion-ids',
               'B2B', 'fulfilled-by', 'Unnamed: 22'],
              dtype='object')
```

```
In [49]: df = df.drop(columns=['Unnamed: 22', 'index'])
```

Rename columns

Let's clean up some column names, ideally in python no spaces. Let's make it all very consistent.

```
In [50]: # Rename confusing column names for clarity
df = df.rename(columns={
    'Order ID': 'order_id',
    'Date': 'date',
    'Status': 'status',
    'Fulfilment': 'fulfilment',
    'Sales Channel ': 'sales_channel',
    'ship-service-level': 'ship_service_level',
    'Style': 'style',
    'Category': 'category',
    'Size': 'size',
    'Courier Status': 'courier_status',
    'Qty': 'quantity',
    'currency': 'currency',
    'Amount': 'amount',
    'ship-city': 'ship_city',
    'ship-state': 'ship_state',
    'ship-postal-code': 'ship_postal_code',
    'ship-country': 'ship_country',
```

```
'promotion-ids': 'promotion_ids',
'fulfilled-by': 'fulfilled_by',
})
df.head(3)
```

Out[50]:

	order_id	date	status	fulfilment	sales_channel	ship_service_level	style	base-SKU	SKU	category	...	quantit
0	8078784-5731545	405-04-30-22	Cancelled	Merchant	Amazon.in	Standard	SET389	SET389	SET389-KR-NP-S	Set	...	
1	9198151-1101146	171-04-30-22	Shipped - Delivered to Buyer	Merchant	Amazon.in	Standard	JNE3781	JNE3781	JNE3781-KR-XXXL	kurta	...	
2	0687676-7273146	404-04-30-22	Shipped	Amazon	Amazon.in	Expedited	JNE3371	JNE3371	JNE3371-KR-XL	kurta	...	

3 rows × 23 columns



Handle missing values

Are there missing values?

In [51]: `df.isna().sum()`

```
Out[51]: order_id      0
date      0
status    0
fulfilment 0
sales_channel 0
ship_service_level 0
style      0
base-SKU   0
SKU        0
category   0
size       0
ASIN       0
courier_status 6866
quantity   0
currency   7786
amount     7786
ship_city   33
ship_state  33
ship_postal_code 33
ship_country 33
promotion_ids 47530
B2B         0
fulfilled_by 88023
dtype: int64
```

Note that `df.isna()` is a data frame of the same shape as `df`, whose terms have type `bool`. With `sum`, we get the number of `True` values per column.

Missing values are denoted by `NaN` in Pandas. When a Pandas object is built, both plain Python's `None` and NumPy's `nan` are taken as `NaN`. Since `np.nan` has type `float`, a numeric series containing `NaN` values gets type `float`.

Three useful Pandas methods related to missing values, which can be applied to both series and data frames, are:

- `isna` returns a Boolean mask indicating which terms are missing.
- `fillna` is used for replacing `NaN`'s by a fixed value, set by the user.
- `dropna` returns the same data frame minus the rows that contain at least one missing value. If a list of columns is specified, the missing values are searched only for those columns.

For our case, let's apply the following business rules:

- If `amount` or `currency` is missing, drop the row (we can't use it for financial analysis).
- If `courier status` is missing, fill it with `"Unknown"`.

The rest we leave as is.

```
In [52]: df = df.dropna(subset=['amount', 'currency']).copy()
df['courier_status'] = df['courier_status'].fillna('Unknown')
```

Remove Duplicates

There are two useful Pandas functions for managing duplicates:

- `drop_duplicates` drops the duplicated entries (in a series) or the duplicated rows (in a data frame).
- `duplicated` returns a Boolean series indicating which entries (for a series) or which rows (for a data frame) are duplicated.

```
In [53]: df.duplicated().sum()
```

```
Out[53]: 3
```

```
In [54]: print('Before dropping duplicates:', df.shape)
df = df.drop_duplicates()
print('After dropping duplicates:', df.shape)
```

Before dropping duplicates: (118728, 23)

After dropping duplicates: (118725, 23)

Any other final conversion

Since I am not too familiar (intuitively) with INR, I prefer to convert to Euro. We could use an API for this conversion, but for the purposes of today an approximate exchange rate suffices.

```
In [55]: exchange_rate = 0.0118 # In 2022
df['amount_eu'] = (df['amount'] * exchange_rate).round(2)
df['unit_price'] = df['amount_eu']
df['revenue'] = df['quantity'] * df['amount_eu']
```

It's good practice to save the dataframe after preprocessing this allows you to continue with a clean dataset going forward.

Make sure you specify a path that makes sense to you. When in doubt you can just use `'./Amazon-cleaned.csv'`

```
In [ ]: df.to_csv('../resources/tabular/Amazon-cleaned.csv', index=False)
```

```
df.to_exce
```

Exploratory Analysis

```
In [2]: import pandas as pd
df = pd.read_csv('../resources/tabular/Amazon-cleaned.csv')
df.head(3)
```

```
Out[2]:
```

	order_id	date	status	fulfilment	sales_channel	ship_service_level	style	base-SKU	SKU	category	...	ship
0	405-8078784-5731545	04-30-22	Cancelled	Merchant	Amazon.in	Standard	SET389	SET389	SET389-KR-NP-S	Set	...	MUN
1	171-9198151-1101146	04-30-22	Shipped - Delivered to Buyer	Merchant	Amazon.in	Standard	JNE3781	JNE3781	JNE3781-KR-XXXL	kurta	...	BENGAL
2	404-0687676-7273146	04-30-22	Shipped	Amazon	Amazon.in	Expedited	JNE3371	JNE3371	JNE3371-KR-XL	kurta	...	MUN

3 rows × 26 columns



Revenue Trend over Time

```
In [3]: df['month'] = pd.to_datetime(df.date).dt.month_name()
df
```

/tmp/ipykernel_88138/785575918.py:1: UserWarning: Could not infer format, so each element will be parsed individually, falling back to `dateutil`. To ensure parsing is consistent and as-expected, please specify a format.

```
df['month'] = pd.to_datetime(df.date).dt.month_name()
```

Out[3]:

	order_id	date	status	fulfilment	sales_channel	ship_service_level	style	base-SKU	SKU	category	...
	0	405-8078784-5731545	04-30-22	Cancelled	Merchant	Amazon.in	Standard	SET389	SET389	SET389-KR-NP-S	Set ... M
	1	171-9198151-1101146	04-30-22	Shipped - Delivered to Buyer	Merchant	Amazon.in	Standard	JNE3781	JNE3781	JNE3781-KR-XXXL	kurta ...
	2	404-0687676-7273146	04-30-22	Shipped	Amazon	Amazon.in	Expedited	JNE3371	JNE3371	JNE3371-KR-XL	kurta ... M
	3	403-9615377-8133951	04-30-22	Cancelled	Merchant	Amazon.in	Standard	J0341	J0341	J0341-DR-L	Western Dress ...
	4	407-1069790-7240320	04-30-22	Shipped	Amazon	Amazon.in	Expedited	JNE3671	JNE3671	JNE3671-TU-XXXL	Top ...
	...	...	...	...	...	...	...	...	...	...	...
	118720	406-6001380-7673107	05-31-22	Shipped	Amazon	Amazon.in	Expedited	JNE3697	JNE3697	JNE3697-KR-XL	kurta ...
	118721	402-9551604-7544318	05-31-22	Shipped	Amazon	Amazon.in	Expedited	SET401	SET401	SET401-KR-NP-M	Set ...
	118722	407-9547469-3152358	05-31-22	Shipped	Amazon	Amazon.in	Expedited	J0157	J0157	J0157-DR-XXL	Western Dress ...
	118723	402-6184140-0545956	05-31-22	Shipped	Amazon	Amazon.in	Expedited	J0012	J0012	J0012-SKD-XS	Set ...
	118724	408-7436540-8728312	05-31-22	Shipped	Amazon	Amazon.in	Expedited	J0003	J0003	J0003-SET-S	Set ... Ct

118725 rows × 27 columns



Adding time

Pandas has powerful tools to help you work with dates - but dates start as plain text ("strings"), which can't be analyzed or sorted properly.

So first, we **convert the date column** into a proper **datetime object**. Once we do that, we unlock a *lot* of cool tricks. Here's the flow:


Concept	What it Does
<code>pd.to_datetime()</code>	Converts text into a proper date format
<code>.dt</code>	A special "datetime accessor" in pandas - lets you pull out month, day, etc.
<code>.dt.month_name()</code>	Extracts the month name (e.g., "April") from a full date

```
In [4]: df['date'] = pd.to_datetime(df['date'])
```

```
/tmp/ipykernel_88138/3532345252.py:1: UserWarning: Could not infer format, so each element will be parsed individually, falling back to `dateutil`. To ensure parsing is consistent and as-expected, please specify a format.  
df['date'] = pd.to_datetime(df['date'])
```

```
In [5]: df['month'] = df.date.dt.month_name()
```

To report sales per month we need to do a `groupby()` (explained in class). This is what it looks like conceptually in our case:



This method groups the rows according to one or more **grouping variables**. Here, you specify `groupby(by='room_type')` .

Next, you select the columns to be aggregated, which, here would be just `price` . Finally, you specify the **aggregation function**. Having only the mean, you can do it just by adding `mean()` or by specifying `.agg('mean')` . The method `round()` works as you would expect it to.

```
In [17]: monthly_revenue = df.groupby('month')[['revenue']].sum()
monthly_revenue
```

```
Out[17]:
```

	revenue
month	
April	328611.70
June	268579.04
May	298847.94

```
In [18]: # Reorder months chronologically (groupby returns them alphabetically)
month_order = ['April', 'May', 'June']
monthly_revenue = monthly_revenue.loc[month_order]
monthly_revenue
```

```
Out[18]:
```

	revenue
month	
April	328611.70
May	298847.94
June	268579.04

```
In [19]: # Calculate average revenue over last 3 months
# and how much below/above average we are
avg_q2_revenue = monthly_revenue['revenue'].mean()
avg_q2_revenue

monthly_revenue['% below avg'] = round((1 - monthly_revenue['revenue'] / avg_q2_revenue) * 100, 1)
monthly_revenue
```

```
Out[19]:
```

	revenue	% below avg
month		
April	328611.70	-10.0
May	298847.94	-0.1
June	268579.04	10.1

Charts are often more persuasive than tables. We will not waste time teaching this as you can ask LLM's to provide this to you. You can "vibe"-chart your way to great visualizations, just make sure that your dataframe above is 100% correct!

```
In [20]: import matplotlib.pyplot as plt
import seaborn as sns

sns.set_style('whitegrid')
fig, ax = plt.subplots(figsize=(8, 6))

# Set-up data
monthly_revenue = df.groupby('month')[['revenue']].sum()
```

```

monthly_revenue = monthly_revenue.T[['April', 'May', 'June']].T
avg_q2_revenue = monthly_revenue['revenue'].mean()
monthly_revenue['% below avg'] = round((1 - monthly_revenue['revenue'] / avg_q2_revenue) * 100, 1)

# Plot bars
bars = ax.bar(monthly_revenue.index, monthly_revenue['revenue'], color='#878787')

# Add horizontal line for average
ax.axhline(avg_q2_revenue, linestyle='--', color='orange', linewidth=2, label='Q2 Average Revenue')

# Annotate each bar
for bar, pct in zip(bars, monthly_revenue['% below avg']):
    ax.annotate(f"{pct}% below avg." if pct > 0 else f"{abs(pct)}% above avg.",
                xy=(bar.get_x() + bar.get_width()/2, bar.get_height() - 7000),
                xytext=(0, 8), textcoords='offset points',
                ha='center', fontsize=14, fontweight='bold')

# Titles and labels
ax.set_title("Amazon India Net Revenue", fontsize=20, x=0.19, y=1.05)
ax.text(-0.08, 1.02, "Q2 FY22", fontsize=15, color='#878787', transform=ax.transAxes)
ax.set_ylabel('Net Revenue in €10,000', fontsize=12, labelpad=3)
ax.set_xlabel(None)

# Ticks and grid
ax.set_yticklabels(list(range(0, 41, 5)))
ax.yaxis.grid(linestyle='--', color='gray', linewidth=0.5, dashes=(8, 5))
ax.xaxis.grid(False)
ax.tick_params(axis='both', labelsize=12)


# Spines
for spine in ['top', 'right', 'left']:
    ax.spines[spine].set_visible(False)
ax.spines['bottom'].set_linewidth(2)
ax.spines['bottom'].set_color('black')

plt.legend(bbox_to_anchor=(1,1.05), fontsize=12, fancybox=True)
plt.show()

```

/tmp/ipykernel_88138/1770982019.py:34: UserWarning: set_ticklabels() should only be used with a fixed number of ticks, i.e. after set_ticks() or using a FixedLocator.
ax.set_yticklabels(list(range(0, 41, 5)))

Amazon India Net Revenue



Revenue declines steadily from April through June. One possible explanation is the Indian **festival calendar**:


Akshaya Tritiya (Apr 3) and Ram Navami (Apr 10) are major auspicious shopping occasions, while Eid al-Fitr (May 2–3; a significant fashion occasion for India's Muslim community (~15% of the population)) likely contributes to the early May uptick. June then quiets as peak summer heat reduces discretionary activity.

You try it:

On which days is Amazon earning most of its revenue in India?

1. First create a column containing the day of the week. Hint: use `pd['<datefield>'].dt.day_name()` for this (day of the week) or alternatively `...dt.day` (day of month).
2. Then, use a groupby and an aggregation function to obtain the total revenue.
3. BONUS: plot the result using `df.plot.bar()` or `df.plot()`


You may want to produce something like this:



```
In [31]: df['day_of_week'] = df['date'].dt.day_name()

df.groupby('day_of_week')\
  [['revenue', 'quantity']]\
  .sum()\
  .plot.bar()
```

Out[31]: <Axes: xlabel='day_of_week'>




```
In [32]: import matplotlib.pyplot as plt

df['day_of_week'] = df['date'].dt.day_name()
day_order = ['Monday', 'Tuesday', 'Wednesday', 'Thursday', 'Friday', 'Saturday', 'Sunday']
revenue_by_day = df.groupby('day_of_week')[['revenue']].sum().reindex(day_order)

revenue_by_day.plot.bar(figsize=(8, 4), color='steelblue')
plt.title('Revenue by Day of Week')
plt.xlabel('')
plt.ylabel('Revenue (€)')
plt.xticks(rotation=45)
```

```
plt.tight_layout()
plt.show()
```



Product Portfolio Analysis

Analysis: revenue across product categories?

→ This feeds Decision 1: which category to invest in next season.

```
In [37]: df['category'].unique()
```

```
Out[37]: array(['Set', 'kurta', 'Western Dress', 'Top', 'Ethnic Dress', 'Bottom',
        'Saree', 'Blouse', 'Dupatta'], dtype=object)
```

Next question: **Which product categories generate the most revenue?**

Typical question to ask:

- Which category dominates revenue? Do we risk of over-dependence?
- Is there a drop-off after top 3 categories or any other long tail effect?
- Are high-volume categories also high-value? Might lead into later analysis (avg. order value).

We'll use `groupby()` again, but this time by `category`. It works just like with months, but this helps us understand which types of products are performing best.

```
In [41]: category_revenue = df.groupby('category')[['revenue']].sum()
category_revenue = category_revenue / category_revenue.sum()
category_revenue
```

```
Out[41]:
```


	revenue
category	
Blouse	0.005808
Bottom	0.001881
Dupatta	0.000012
Ethnic Dress	0.010033
Saree	0.001656
Set	0.498827
Top	0.068983
Western Dress	0.140945
kurta	0.271854

```
In [42]: category_revenue = df.groupby('category')[['revenue']].sum()
category_revenue = category_revenue / category_revenue.sum() * 100
category_revenue = category_revenue.sort_values('revenue', ascending=False)

category_revenue.plot.bar()
```

```
plt.title('Revenue by Product Category (%)')
plt.ylabel('% of Total Revenue')
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()
```

category_revenue



Out[42]:

revenue	
category	
Set	49.882695
kurta	27.185384
Western Dress	14.094520
Top	6.898323
Ethnic Dress	1.003319
Blouse	0.580821
Bottom	0.188110
Saree	0.165622
Dupatta	0.001205

For reference, here are the referred to styles (Kurta - right, Western dress - left):



Analysis: Average Order Value by Category


Revenue alone doesn't tell the full story. Some categories might generate fewer sales but have *higher order values*.

Let's calculate **average revenue per order** (AOV) for each product category.

You try it


Tabulate and then plot the average order value per category.

You may want to produce something like this:




```
In [44]: df.groupby('category')[['revenue']].mean().plot.bar()
```

```
Out[44]: <Axes: xlabel='category'>
```



```
In [45]: category_aov = df.groupby('category')['revenue'].mean().sort_values(ascending=False)
category_aov.plot.bar(color='coral')
plt.title('Average Order Value by Category')
plt.ylabel('AOV (€)')
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()

category_aov
```



```
Out[45]: category
Set          9.720927
Saree        9.636623
Western Dress 8.723655
Ethnic Dress  8.481255
Top           6.188590
Blouse        6.065711
kurta         5.319058
Bottom        4.121125
Dupatta       3.600000
Name: revenue, dtype: float64
```

Let's add line revenue for contrast

```
In [47]: fig, ax = plt.subplots(figsize=(10,6))

bar_colors = ['green' if cat in ['Western Dress', 'Set'] else 'coral' for cat in category_revenue.index]


category_revenue.plot(kind='bar', color='coral', ax=ax)

for i, category in enumerate(category_revenue.index):
    avg_cost = df[df['category'] == category]['revenue'].mean()
    ax.text(i, category_revenue.values[i] + .1, f'€{avg_cost:.2f}', ha='center', fontsize=10)

ax.annotate(
    'High Potential?',
    xy=(2, 15),
    xytext=(2, 15 + 5),
    arrowprops=dict(facecolor='black', shrink=0.15, width=2, headwidth=8),
    ha='center',
    fontsize=11,
    fontweight='bold'
)

ax.annotate(
    'High Potential?',
    xy=(4, 2),
    xytext=(2, 20),
    arrowprops=dict(facecolor='black', shrink=0.05, width=2, headwidth=8),
    ha='center',
    fontsize=11,
    fontweight='bold'
)

plt.title('Revenue by Product Category', fontsize=14)
plt.xlabel('Product Category')
plt.ylabel('Percentage of Total Revenue')
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()
```

Analysis: Product Portfolio Matrix

We can think of each product category as a portfolio element. The right lens here is **Revenue = Volume × AOV** — these are the two independent levers you can pull. Plotting volume (order count) against AOV puts each category in a strategic quadrant:

Quadrant	Lever to pull	Action
Volume Drivers	↑ AOV	Bundle, upsell, larger pack sizes — the volume is there, now raise the basket
Niche Upside	↑ Volume	Paid visibility, promos, SEO — AOV justifies acquisition spend
Premium Performers	Protect both	Defend share, invest in retention, candidate for expansion
Underperformers	Neither working	Investigate root cause before retiring — could be niche or seasonal

```
In [53]: import matplotlib.pyplot as plt

# Calculate total revenue and AOV per category
category_metrics = df.groupby('SKU').agg(
    total_revenue = ('revenue', 'sum'),
    avg_order_value = ('revenue', 'mean')
).reset_index()
category_metrics
```

```
Out[53]:
```

	SKU	total_revenue	avg_order_value
0	AN201-RED-M	5.40	2.700
1	AN201-RED-XL	7.10	3.550
2	AN201-RED-XXL	2.70	2.700
3	AN202-ORANGE-M	6.25	3.125
4	AN202-ORANGE-S	13.46	3.365
...	...	...	...
7134	SET449-KR-NP-S	17.10	17.100
7135	SET461-KR-NP-XL	12.26	12.260
7136	SET461-KR-NP-XS	9.07	9.070
7137	SET461-KR-NP-XXL	9.07	9.070
7138	SET462-KR-NP-S	11.79	11.790

7139 rows × 3 columns

You try it

Try to scatter plot the total revenue v.s. median order value. The command you'll use is `df.plot.scatter(x=...,y=...)`.


```
In [52]: category_metrics = df.groupby('category').agg(
        total_revenue = ('revenue', 'sum'),
        avg_order_value = ('revenue', 'mean'),
    ).reset_index()
category_metrics['revenue_10k'] = category_metrics['total_revenue'] / 10000

fig, ax = plt.subplots(figsize=(9, 6))
ax.scatter(category_metrics['revenue_10k'], category_metrics['avg_order_value'], s=200, alpha=0.7)

for _, row in category_metrics.iterrows():
    ax.text(row['revenue_10k'], row['avg_order_value'], row['category'], fontsize=9, ha='right')

ax.axhline(category_metrics['avg_order_value'].median(), color='gray', linestyle='--', linewidth=1)
ax.axvline(category_metrics['revenue_10k'].median(), color='gray', linestyle='--', linewidth=1)

ax.set(title='Product Categories: Revenue vs. AOV',
        xlabel='Total Revenue (€10,000)', ylabel='Average Order Value (€)')
ax.grid(True, linestyle='--', alpha=0.5)
plt.tight_layout()
plt.show()
```



```
In [50]: # --- filler: ensures notebook runs ---
category_metrics = df.groupby('category')['revenue'].agg(
    total_revenue='sum',
    avg_order_value='mean'
).reset_index()
```

```
In [51]: portfolio = df.groupby('category').agg(
    order_count = ('order_id', 'count'),
    aov = ('revenue', 'mean'),
    revenue = ('revenue', 'sum')
).reset_index()

med_volume = portfolio['order_count'].median()
med_aov = portfolio['aov'].median()

x_min, x_max = 0, portfolio['order_count'].max() * 1.15
y_min, y_max = portfolio['aov'].min() * 0.85, portfolio['aov'].max() * 1.10

fig, ax = plt.subplots(figsize=(9, 6))

# Shade quadrants
ax.fill_between([x_min, med_volume], y_min, med_aov, color='#e74c3c', alpha=0.08)
```



```

ax.fill_between([med_volume, x_max], y_min, med_aov, color='#f39c12', alpha=0.08)
ax.fill_between([x_min, med_volume], med_aov, y_max, color='#3498db', alpha=0.08)
ax.fill_between([med_volume, x_max], med_aov, y_max, color='#2ecc71', alpha=0.10)

ax.axvline(med_volume, color='gray', linestyle='--', linewidth=1)
ax.axhline(med_aov, color='gray', linestyle='--', linewidth=1)


sizes = (portfolio['revenue'] / portfolio['revenue'].max() * 2000).clip(lower=80)
ax.scatter(portfolio['order_count'], portfolio['aov'], s=sizes, alpha=0.7, color='#333333', zorder=3)

for _, row in portfolio.iterrows():
    ax.text(row['order_count'], row['aov'] + 0.15, row['category'], fontsize=9, ha='center', zorder=4)

t = ax.transAxes
ax.text(0.02, 0.98, 'Niche Upside', transform=t, va='top', ha='left', fontsize=9, color='#2980b9',
ax.text(0.98, 0.98, 'Premium Performers', transform=t, va='top', ha='right', fontsize=9, color='#27ae60',
ax.text(0.02, 0.02, 'Underperformers', transform=t, va='bottom', ha='left', fontsize=9, color='#c0392b',
ax.text(0.98, 0.02, 'Volume Drivers', transform=t, va='bottom', ha='right', fontsize=9, color='#e67e22',

ax.set_xlim(x_min, x_max)
ax.set_ylim(y_min, y_max)
ax.set(title='Product Portfolio: Volume vs AOV (bubble = revenue)',
       xlabel='Order Count (volume)', ylabel='Average Order Value (€)')
ax.grid(False)
plt.tight_layout()
plt.show()

```



Here are the two interesting ones that come up:



Cancellations & Returns

Not all orders are good news. In e-commerce, cancellations and returns create serious operational and financial issues:

- Lost revenue: These orders don't convert into sales.
- Shipping and handling costs: Often non-refundable.
- Reverse logistics: Costly and time-consuming.
- Chargebacks: Especially if customers dispute transactions.
- Fraud & abuse: Repeated returns or B2C-to-B2B loopholes.
- Money laundering flags: In some cases, abnormal return behavior is a red flag.

Long story short: cancellations and returns hurt revenue.

Let's measure the size of the problem:


- How many orders were cancelled?
- How many were returned?
- What percentage of total orders does this represent?

So far, you've used `.groupby()` for totals and averages. This time let's introduce `.value_counts()` which is a quick way to count categories in one column, like a faster `.groupby('status').shape`.

→ *This feeds Decision 2: where to focus the operational fix.*

I start by thinking how I will visualize this. I'm usually not a fan of pie charts (hard to cognitively process proportions from angles), but given that we have so many barcharts on the dashboard I think it will add some visual appeal.

Maybe something like this in the Amazon colors?



Let's go!

```
In [80]: # Step 1: Count all status values
# ---> same as: df.groupby('status').size().sort_values(ascending=False)
print(df['status'].value_counts() )

# Step 2: Filter for problematic orders (cancelled or returned)
cancelled = df[df['status'] == 'Cancelled']
returned  = df[df['status'] == 'Shipped - Returned to Seller']


# Step 3: Calculate totals and percentages
total_orders = df.shape[0]
cancelled_pct = cancelled.shape[0] / total_orders * 100
returned_pct  = returned.shape[0] / total_orders * 100
problematic_pct = (cancelled.shape[0] + returned.shape[0]) / total_orders * 100

print("=====")
print(f"Cancelled Orders: {len(cancelled)} ({cancelled_pct:.2f}%)")
print(f"Returned Orders: {len(returned)} ({returned_pct:.2f}%)")
print(f"Total Problematic Orders: {problematic_pct:.2f}%")

print("=====")
plt.pie(
    [len(cancelled), len(returned), len(cancelled) + len(returned)],
    labels=['Cancelled', 'Returned', 'Problematic'],
    colors=['#FFD08A', '#FFB244', '#FF9900'],
    autopct='%1.1f%%'
)
plt.show()
```

```
status
Shipped                75940
Shipped - Delivered to Buyer    28029
Cancelled              10756
Shipped - Returned to Seller    1949
Shipped - Picked Up           945
Pending                 639
Pending - Waiting for Pick Up   272
Shipped - Returning to Seller   143
Shipped - Out for Delivery      35
Shipped - Rejected by Buyer     11
Shipped - Lost in Transit       5
Shipped - Damaged              1
Name: count, dtype: int64

====
Cancelled Orders: 10756 (9.06%)
Returned Orders: 1949 (1.64%)
Total Problematic Orders: 10.70%
=====
```




State Dashboard

The goal is often not just to analyze but rather to communicate insights clearly as we've been trying to do before. Sometimes a plot is enough, other times a dynamic dashboard is needed.

Dashboards allow other users (e.g., regional managers) to quickly:

- See how their region is doing
- Spot product trends
- Act on poor performance
- ..

We'll build a simple command-line dashboard first, and then make it interactive. We're working towards something like this that shows per-state statistics:



Let's first look at the states:

→ This gives us situational awareness: where is our business concentrated?

```
In [65]: df.ship_state.unique()
```

```
Out[65]: array(['MAHARASHTRA', 'KARNATAKA', 'PUDUCHERRY', 'TAMIL NADU',
                'UTTAR PRADESH', 'CHANDIGARH', 'TELANGANA', 'ANDHRA PRADESH',
                'RAJASTHAN', 'DELHI', 'HARYANA', 'ASSAM', 'JHARKHAND',
                'CHHATTISGARH', 'ODISHA', 'KERALA', 'MADHYA PRADESH',
                'WEST BENGAL', 'NAGALAND', 'GUJARAT', 'UTTARAKHAND', 'BIHAR',
                'JAMMU & KASHMIR', 'PUNJAB', 'HIMACHAL PRADESH',
                'ARUNACHAL PRADESH', 'MANIPUR', 'GOA', 'MEGHALAYA', 'TRIPURA',
                'LADAKH', 'DADRA AND NAGAR', 'SIKKIM', nan, 'ANDAMAN & NICOBAR',
                'RAJSTHAN', 'NL', 'MIZORAM', 'NEW DELHI',
                'PUNJAB/MOHALI/ZIRAKPUR', 'RJ', 'ORISSA', 'LAKSHADWEEP', 'PB',
                'AR', 'PONDICHERRY', 'RAJSTHAN'], dtype=object)
```

Let's pick just one state to understand its sales profile. We'll start with Karnataka, one of the largest markets

```
In [66]: # Filter data for Karnataka
karnataka = df[df['ship_state'] == 'KARNATAKA']
```

```
In [67]: # Calculate the number of orders placed in this state
total_orders = len(karnataka)

# Calculate total revenue generated
total_revenue = karnataka['amount_eu'].sum()

# Average revenue per order – helps identify spending behavior
avg_order_value = karnataka['amount_eu'].mean()

# Find the most popular product category by quantity sold
top_category = (
    karnataka.groupby('category')['quantity']
    .sum()
    .sort_values(ascending=False)
    .head(1)
)

# Print the summary stats
print(f"Total Orders: {total_orders}")
print(f"Total Revenue: €{total_revenue:,.2f}")
print(f"Average Order Value: €{avg_order_value:,.2f}")
print(f"Top Category: {top_category.index[0]} ({top_category.values[0]} units)")
```

Total Orders: 16050
 Total Revenue: €123,519.42
 Average Order Value: €7.70
 Top Category: kurta (6085 units)

You try it (at home): write a dashboard function `draw_state_dashboard(state)` that takes as input a `state` and prints the above statistics.

```
In [68]: def draw_state_dashboard(state):
    subset = df[df['ship_state'] == state.upper()]
    if subset.empty:
        print(f"No data for {state}")
        return

    total_orders = len(subset)
    total_revenue = subset['revenue'].sum()
    aov = subset['revenue'].mean()
    top_category = subset.groupby('category')['revenue'].sum().idxmax()

    print(f"--- {state.title()} ---")
    print(f"  Orders:      {total_orders:,}")
    print(f"  Total Revenue: €{total_revenue:,.0f}")
    print(f"  AOV:          €{aov:.2f}")
    print(f"  Top Category: {top_category}")

draw_state_dashboard('KARNATAKA')
```

```
--- Karnataka ---
Orders:      16,050
Total Revenue: €120,478
AOV:          €7.51
Top Category: Set
```

Dashboard

```
In [69]: # --- filler: ensures notebook runs ---
def draw_state_dashboard(state):
    state_df = df[df['ship_state'] == state]
    orders = len(state_df)
    revenue = state_df['revenue'].sum()
    aov = state_df['revenue'].mean()
    top_cat = state_df.groupby('category')['revenue'].sum().idxmax()
    print(f"State: {state}")
    print(f"  Orders:      {orders:,}")
    print(f"  Total Revenue: €{revenue:,.0f}")
    print(f"  AOV:          €{aov:.2f}")
    print(f"  Top Category: {top_cat}")
```

```
In [70]: # See solution at end of notebook
```

BONUS: add some analysis from before to make it prettier

Note: in production you will want to use `plotly dash`, which functions similar to what we've seen above. Here are a [couple of examples](#), and one for [Amazon](#).

In [resources](#) you'll find the full code to make an interactive dashboard based on the above functions using python.

Solutions

Solution: Category Revenue Analysis


```
In [81]: category_revenue = df.groupby('category')[['revenue']].sum()
category_revenue = category_revenue / category_revenue.sum() * 100

category_revenue = category_revenue.sort_values(ascending=False, by='revenue')
category_revenue
```

```
Out[81]:
```

	revenue
category	
Set	49.882695
kurta	27.185384
Western Dress	14.094520
Top	6.898323
Ethnic Dress	1.003319
Blouse	0.580821
Bottom	0.188110
Saree	0.165622
Dupatta	0.001205

```
In [82]: # Basic plot command
category_revenue.plot.bar();
```



Bonus

Revenue by Customer Type (B2B vs. Regular)

Amazon doesn't just sell to everyday consumers — it also serves **business buyers** through **Amazon Business**.

Customer Type	Examples	What They Value
Regular (B2C)	Individual shoppers buying for personal use	Price, fast shipping, product variety
Business (B2B)	SMEs, schools, retailers buying in bulk	Discounts, tax invoices, delivery reliability, APIs

Business customers usually spend more per order — but place fewer orders.

Regular customers are more frequent, but each transaction is smaller.

To understand these channels better, let us look at some stats:

1. Total revenue by customer type
2. Average order value (AOV) by customer type
3. Number of orders by customer type

In [73]: *# Step 1: Group by customer type and summarize metrics*

```
customer_summary = df.groupby('B2B').agg(
    total_revenue = ('amount_eu', 'sum'),
    avg_order_value = ('amount_eu', 'mean'),
    order_count = ('order_id', 'count')
).reset_index()
customer_summary
```

Out[73]:

	B2B	total_revenue	avg_order_value	order_count
0	False	919207.39	7.797691	117882
1	True	6976.65	8.275979	843

0	False	919207.39	7.797691	117882
1	True	6976.65	8.275979	843

In [74]: *# Step 2: Convert True/False to readable labels*

```
customer_summary['B2B'] = customer_summary['B2B'].replace({True: 'B2B', False: 'B2C'})
customer_summary
```

Out[74]:

	B2B	total_revenue	avg_order_value	order_count
0	B2C	919207.39	7.797691	117882
1	B2B	6976.65	8.275979	843

0	B2C	919207.39	7.797691	117882
1	B2B	6976.65	8.275979	843

Observations:

Metric	What It Shows
Order Count	Regular customers dominate volume — this is the bulk of the business
AOV	Business customers spend ~8% more per order (€7.71 vs €7.13) but much less than expected.
Total Revenue	Business is currently a <i>tiny slice</i> of total revenue (~0.75%) — but with higher ticket size

Potential actionable insight: one last attempt at invigorating and attracting better B2B through e.g. exclusive deals, bulk pricing, repeat order flows? If it doesn't work we may consider spinning off or integrating it into B2C.

Interactive State Selector (ipywidgets)

The dashboard in `resources/amazon/dashboard/` is a production-quality version of this. The cell below shows the interactive widget version for use directly in Jupyter.

Wonderful, all that is needed at this point is to make it interactive with a dropdown. Note: this code assumes that you are using a notebook like google collab or jupyter lab. Notably it won't work in VS Code.

Note: You will need to install and enable widgets, it doesn't always work on every environment.

In [75]:

```
! pip install ipywidgets==8.0.2
! pip install jupyterlab_widgets==3.0.13
```


Requirement already satisfied: ipywidgets==8.0.2 in /home/enric/miniconda3/envs/dsfb/lib/python3.11/site-packages (8.0.2)

Requirement already satisfied: ipykernel>=4.5.1 in /home/enric/miniconda3/envs/dsfb/lib/python3.11/site-packages (from ipywidgets==8.0.2) (6.29.4)

Requirement already satisfied: ipython>=6.1.0 in /home/enric/miniconda3/envs/dsfb/lib/python3.11/site-packages (from ipywidgets==8.0.2) (8.24.0)

Requirement already satisfied: traitlets>=4.3.1 in /home/enric/miniconda3/envs/dsfb/lib/python3.11/site-packages (from ipywidgets==8.0.2) (5.14.3)

Requirement already satisfied: widgetsnbextension~=4.0 in /home/enric/miniconda3/envs/dsfb/lib/python3.11/site-packages (from ipywidgets==8.0.2) (4.0.13)

Requirement already satisfied: jupyterlab-widgets~=3.0 in /home/enric/miniconda3/envs/dsfb/lib/python3.11/site-packages (from ipywidgets==8.0.2) (3.0.13)

Requirement already satisfied: comm>=0.1.1 in /home/enric/miniconda3/envs/dsfb/lib/python3.11/site-packages (from ipykernel>=4.5.1->ipywidgets==8.0.2) (0.2.2)

Requirement already satisfied: debugpy>=1.6.5 in /home/enric/miniconda3/envs/dsfb/lib/python3.11/site-packages (from ipykernel>=4.5.1->ipywidgets==8.0.2) (1.8.1)

Requirement already satisfied: jupyter-client>=6.1.12 in /home/enric/miniconda3/envs/dsfb/lib/python3.11/site-packages (from ipykernel>=4.5.1->ipywidgets==8.0.2) (8.6.3)

Requirement already satisfied: jupyter-core!=5.0.*,>=4.12 in /home/enric/miniconda3/envs/dsfb/lib/python3.11/site-packages (from ipykernel>=4.5.1->ipywidgets==8.0.2) (5.7.2)

Requirement already satisfied: matplotlib-inline>=0.1 in /home/enric/miniconda3/envs/dsfb/lib/python3.11/site-packages (from ipykernel>=4.5.1->ipywidgets==8.0.2) (0.1.7)

Requirement already satisfied: nest-asyncio in /home/enric/miniconda3/envs/dsfb/lib/python3.11/site-packages (from ipykernel>=4.5.1->ipywidgets==8.0.2) (1.6.0)

Requirement already satisfied: packaging in /home/enric/miniconda3/envs/dsfb/lib/python3.11/site-packages (from ipykernel>=4.5.1->ipywidgets==8.0.2) (23.2)

Requirement already satisfied: psutil in /home/enric/miniconda3/envs/dsfb/lib/python3.11/site-packages (from ipykernel>=4.5.1->ipywidgets==8.0.2) (5.9.8)

Requirement already satisfied: pyzmq>=24 in /home/enric/miniconda3/envs/dsfb/lib/python3.11/site-packages (from ipykernel>=4.5.1->ipywidgets==8.0.2) (26.0.3)

Requirement already satisfied: tornado>=6.1 in /home/enric/miniconda3/envs/dsfb/lib/python3.11/site-packages (from ipykernel>=4.5.1->ipywidgets==8.0.2) (6.3.3)

Requirement already satisfied: decorator in /home/enric/miniconda3/envs/dsfb/lib/python3.11/site-packages (from ipython>=6.1.0->ipywidgets==8.0.2) (5.1.1)

Requirement already satisfied: jedi>=0.16 in /home/enric/miniconda3/envs/dsfb/lib/python3.11/site-packages (from ipython>=6.1.0->ipywidgets==8.0.2) (0.19.1)

Requirement already satisfied: prompt-toolkit<3.1.0,>=3.0.41 in /home/enric/miniconda3/envs/dsfb/lib/python3.11/site-packages (from ipython>=6.1.0->ipywidgets==8.0.2) (3.0.48)

Requirement already satisfied: pygments>=2.4.0 in /home/enric/miniconda3/envs/dsfb/lib/python3.11/site-packages (from ipython>=6.1.0->ipywidgets==8.0.2) (2.18.0)

Requirement already satisfied: stack-data in /home/enric/miniconda3/envs/dsfb/lib/python3.11/site-packages (from ipython>=6.1.0->ipywidgets==8.0.2) (0.6.2)

Requirement already satisfied: typing-extensions>=4.6 in /home/enric/miniconda3/envs/dsfb/lib/python3.11/site-packages (from ipython>=6.1.0->ipywidgets==8.0.2) (4.13.2)

Requirement already satisfied: pexpect>4.3 in /home/enric/miniconda3/envs/dsfb/lib/python3.11/site-packages (from ipython>=6.1.0->ipywidgets==8.0.2) (4.9.0)

Requirement already satisfied: parso<0.9.0,>=0.8.3 in /home/enric/miniconda3/envs/dsfb/lib/python3.11/site-packages (from jedi>=0.16->ipython>=6.1.0->ipywidgets==8.0.2) (0.8.4)

Requirement already satisfied: python-dateutil>=2.8.2 in /home/enric/miniconda3/envs/dsfb/lib/python3.11/site-packages (from jupyter-client>=6.1.12->ipykernel>=4.5.1->ipywidgets==8.0.2) (2.9.0.post0)

Requirement already satisfied: platformdirs>=2.5 in /home/enric/miniconda3/envs/dsfb/lib/python3.11/site-packages (from jupyter-core!=5.0.*,>=4.12->ipykernel>=4.5.1->ipywidgets==8.0.2) (4.3.7)

Requirement already satisfied: ptyprocess>=0.5 in /home/enric/miniconda3/envs/dsfb/lib/python3.11/site-packages (from pexpect>4.3->ipython>=6.1.0->ipywidgets==8.0.2) (0.7.0)

Requirement already satisfied: wcwidth in /home/enric/miniconda3/envs/dsfb/lib/python3.11/site-packages (from prompt-toolkit<3.1.0,>=3.0.41->ipython>=6.1.0->ipywidgets==8.0.2) (0.2.13)

Requirement already satisfied: executing>=1.2.0 in /home/enric/miniconda3/envs/dsfb/lib/python3.11/site-packages (from stack-data->ipython>=6.1.0->ipywidgets==8.0.2) (2.0.1)

Requirement already satisfied: asttokens>=2.1.0 in /home/enric/miniconda3/envs/dsfb/lib/python3.11/site-packages (from stack-data->ipython>=6.1.0->ipywidgets==8.0.2) (2.4.1)

Requirement already satisfied: pure-eval in /home/enric/miniconda3/envs/dsfb/lib/python3.11/site-packages (from stack-data->ipython>=6.1.0->ipywidgets==8.0.2) (0.2.3)

Requirement already satisfied: six>=1.12.0 in /home/enric/miniconda3/envs/dsfb/lib/python3.11/site-packages (from asttokens>=2.1.0->stack-data->ipython>=6.1.0->ipywidgets==8.0.2) (1.17.0)

Requirement already satisfied: jupyterlab_widgets==3.0.13 in /home/enric/miniconda3/envs/dsfb/lib/python3.11/site-packages (3.0.13)

It's possible that you need to restart your kernel after doing this. If that's the case, do so now and run the cell after this:

```
In [76]: import pandas as pd
        ### Download cleaned data
        # df = pd.read_csv("https://raw.githubusercontent.com/ciri/iese-dsfb/refs/heads/main/resources/tabular/Amazon")
        ### Or use your local copy
        # df = pd.read_csv("../resources/tabular/Amazon-cleaned.csv")
```

```
In [77]: import ipywidgets as widgets
        from IPython.display import display

        output = widgets.Output()
        state_dropdown = widgets.Dropdown(
            options=sorted(df['ship_state'].dropna().unique()),
            description='Select a State:',
```



```
        style={'description_width': 'initial'}
    )

def on_state_change(change):
    if change['type'] == 'change' and change['name'] == 'value':
        output.clear_output(wait=True)
        with output:
            draw_state_dashboard(change['new'])

state_dropdown.observe(on_state_change)
display(state_dropdown, output);
```

Dropdown(description='Select a State:', options=('ANDAMAN & NICOBAR', 'ANDHRA PRADESH', 'AR', 'ARUNACHAL PRAD
E...
Output()