

Airline Text Analytics (*solved*)

Airline companies are in fierce competition. One key dimension of competition is customer service, a direct driver of customer satisfaction:



As part of an effort to better understand passenger satisfaction within the Star Alliance network, we analyze customer reviews from thousands of airline passengers. The goal is to identify how factors like sentiment, seat type, travel purpose, and specific routes relate to overall recommendation rates.

The dataset includes 23,000+ reviews, with key variables such as:

- **Airline Name** : the name of the airline.
- **Route** : the flight route taken by the passenger.
- **Review** : free-text review written by the traveler.
- **Type Of Traveller** : e.g., Solo Leisure, Business, Couple.
- **Recommended** : whether the traveler recommends the airline (yes/no).
- **Seat Type** : seat class (e.g., Economy, Business).

Data Source: [Airlinequality.com](https://www.airlinequality.com) via [Kaggle](https://www.kaggle.com)

Required packages — run once before starting:

```
pip install pandas matplotlib seaborn nltk tqdm wordcloud transformers torch google-genai
# or with conda:
# conda install pandas matplotlib seaborn nltk tqdm && pip install wordcloud transformers torch google-genai
```

Part I: Working with Text Data

Strings

Text data always starts out in the form of a string. A **string** is a sequence (list) of **characters**. This includes the (English) alphanumeric characters and also special characters like white space, punctuation, etc. Other symbols, like emoticons, can also appear in your data, specially in social networks data. Besides that, you can also find letters from other languages (Spanish, Portuguese, etc) or alphabets (Cyrillic, hiragana, etc), and even ideographs (such as Han characters).

Characters as numbers

There is a basic set of 127 characters, called the **ASCII characters**, which are encoded in the same way by all the computers, so you will never have trouble with them. They include the English letters (without accents), the numbers, basic punctuation (not curly quote marks or long dashes), white space, **control characters** such as the new line, represented in programming languages as `\n`, and other symbols familiar to you, such as the dollar (\$) and the hash (#) symbols. The complete list can be easily found in Internet.

Non-ASCII characters can be encoded by different computers or different text editors in different ways. Mind that, if you capture string data on your own, you will probably find some of these characters in your data. Even when the documents are expected to be in English, they can be contaminated by other languages: Han characters, German dieresis, Spanish eñe, etc.

The preferred **encoding** is **UTF-8 (utf-8)**, which is the default encoding in Macintosh computers. Reading and writing text files in Pandas, the argument encoding allows you to manage both UTF-8 and the alternative encoding **Latin-1 (latin1)**. Windows computers use their own system, which is region specific. In US and Western Europe, this is **Windows-1252**, which is very close to Latin-1, though not exactly the same.

```
In [5]: the_letter_A = chr(65) # this
        the_letter_A = 'A'   # is the same as this

        print(the_letter_A)

        # This one is in Farsi, it may not render on your computer
        print(chr(0x0604))

        # This one is traditional chinese
        print(chr(0x9F8D))
```

Strings as lists

Both lists and strings are **sequences**, so they share some core behaviours. A string can be thought of as a list of characters. You can use `len()` and slicing just like a list:

```
In [6]: school = 'IESE Business School'
print(len(school))
print(school[:4])    # first 4 characters
print(school[-6:])   # last 6 characters
```

```
20
IESE
School
```

Python string methods

Python has a rich set of methods for manipulating strings:

Method	What it does
<code>lower()</code> / <code>upper()</code>	Change case
<code>replace(old, new)</code>	Find and replace
<code>split(sep)</code>	Split into a list of parts
<code>join(list)</code>	Join a list back into a string
<code>count(pattern)</code>	Count occurrences

```
In [7]: name = 'Enric Junque de Fortuny'
print(name.upper())
print(name.lower())
print(name.split(' '))
print('---'.join(name.split(' ')))
```

```
ENRIC JUNQUE DE FORTUNY
enric junque de fortuny
['Enric', 'Junque', 'de', 'Fortuny']
Enric---Junque---de---Fortuny
```

Airline Reviews Dataset

Let's load the data. I've already partially cleaned it up for you (no duplicates). You can download it from [here](#).

```
In [8]: import pandas as pd

pd.set_option('display.max_colwidth', None)

## Use your local copy:
df = pd.read_csv('../resources/text/Airline_Reviews_v2.csv')[['Airline Name', 'From', 'To', 'Via', 'Review', 'Type']]
## Or download every time:
# df = pd.read_csv('https://raw.githubusercontent.com/ciri/iese-dsfb/refs/heads/main/resources/text/Airline_F

df.head(2)
```

Out[8]:	Airline Name	From	To	Via	Review	Type Of Traveller	Recommended	Seat Type
0	AB Aviation	Moroni	Moheli	NaN	Moroni to Moheli. Turned out to be a pretty decent airline. Online booking worked well, checkin and boarding was fine and the plane looked well maintained. Its a very short flight - just 20 minutes or so so i didn't expect much but they still managed to hand our a bottle of water and some biscuits which i though was very nice. Both flights on time.	Solo Leisure	1	Economy Class
1	AB Aviation	Moroni	Anjouan	NaN	Moroni to Anjouan. It is a very small airline. My ticket advised me to turn up at 0800hrs which I did. There was confusion at this small airport. I was then directed to the office of AB Aviation which was still closed. It opened at 0900hrs and I was told that the flight had been put back to 1300hrs and that they had tried to contact me. This could not be true as they did not have my phone number. I was with a local guide and he had not been informed either. I presume that I was bumped off. The later flight did operate but as usual, there was confusion at check-in. The flight was only 30mins and there were no further problems. Not a good airline but it is the only one for Comoros.	Solo Leisure	0	Economy Class

Exploring the dataset

There is not much to `describe()` here since most columns are text. Let's use `.value_counts()` to count categories, and `.groupby()` to aggregate.

```
In [9]: df['Seat Type'].value_counts()
```

```
Out[9]: Seat Type
Economy Class    19129
Business Class    2097
Premium Economy    646
First Class       186
Name: count, dtype: int64
```

```
In [10]: df.groupby('Seat Type').count()
```

```
Out[10]:
```

	Airline Name	From	To	Via	Review	Type Of Traveller	Recommended
Seat Type							
Business Class	2097	1806	1805	517	2097	1808	2097
Economy Class	19129	16747	16731	3533	19129	16829	19129
First Class	186	168	168	23	186	168	186
Premium Economy	646	608	605	145	646	613	646

```
In [11]: df.groupby('Seat Type').nunique()
```

```
Out[11]:
```

	Airline Name	From	To	Via	Review	Type Of Traveller	Recommended
Seat Type							
Business Class	217	436	471	175	2097	4	2
Economy Class	489	2025	2271	677	19129	4	2
First Class	75	103	106	19	186	4	2
Premium Economy	157	240	260	83	646	4	2

```
In [12]: import matplotlib.pyplot as plt
import seaborn as sns


routes = (
    df.dropna(subset=['From', 'To'])
    .groupby(['From', 'To'])
    .size()
    .reset_index(name='count')
)
top_from = routes['From'].value_counts().head(20).index
top_to = routes['To'].value_counts().head(20).index
```

```

sub = routes[
    routes['From'].isin(top_from) &
    routes['To'].isin(top_to)
]
matrix = sub.pivot(index='From', columns='To', values='count').fillna(0)

plt.figure(figsize=(10,8))
sns.heatmap(matrix, cmap='Blues', vmin=0.0001)
plt.title('Route Intensity Heatmap')
plt.xlabel('To')
plt.ylabel('From')
plt.tight_layout()
plt.show()

```




```

In [13]: top_via = df['Via'].value_counts().head(10)

plt.figure(figsize=(6,4))
sns.barplot(x=top_via.values, y=top_via.index)
plt.title('Most Common Via Stops')
plt.xlabel('Count')
plt.ylabel('Via')
plt.tight_layout()
plt.show()

```



Exploring the review text

When a Pandas column contains string, we can access its **string methods**. Pandas string methods are vectorized versions of the above methods. They return a series of the same length, in which each term results from applying a string function to the corresponding term of `s`. The syntax of these functions is typically `s.str.fname(args)`, where `s` is a Pandas series of strings (or a list of strings) and `args` is a set of ad hoc arguments.

For example, `.str.contains("pattern")` returns a Boolean series — True wherever the pattern appears:

```
In [14]: f_has_space = df.Review.str.lower().str.contains("space")
df[f_has_space].head(3)
```

Out[14]:

	Airline Name	From	To	Via	Review	Type Of Traveller	Recommended	Seat Type
41	Adria Airways	Manchester	Ljubljana	NaN	Excellent flights both ways with Adria Airways from Manchester to Ljubljana. Staff friendly but unobtrusive. Leg room far better than most economy flights, seats comfy and plenty of storage space. Flights on time, quick and efficient boarding and disembarking, smooth take offs and landings. All in all a very pleasant experience.	Family Leisure	1	Economy Class
100	Aegean Airlines	Athens	Chania	NaN	We took 3 flights with Aegean and would not recommend them to anyone. Had 3 issues: 1. They overbooked our flight and bumped many people. We bought our flight 3 months ahead, but checked in about 18 hours before our flight (they allow check-in up to 24 hours ahead). We were issued a "standby" ticket without any seats. When we got to the gate, we were told we may or may not get onto the flight and if we didn't then we would have to wait 4 hours for the late night flight, or try again in the morning. Lots of other people in the same situation. We were never issued boarding passes but we kept following up with the person at the ticket counter until they gave us passes, it seems like due to our persistence and nothing systematic. Luckily we got on that flight, but others didn't. Nowhere in the policy does it say that whether or not you get on the flight is related to your check-in time. If you take Aegean, check in right away and be persistent about asking for seats! 2. They apply a checked bag policy at will. We were sitting near the gate. The gate attendants approached us and asked to weigh our bags. Our bags, which fit easily into any overhead bin (including previous Aegean flights) were found to be slightly overweight (apparently 9kgs are allowed, ours were about 10kgs). We asked if we could move 1kg into our purse / backpack. They said no, they need to check the two bags, and by the way they are charging us 40Euro per bag. This was for a 45 min flight from an island to Athens. We had no choice and were forced to pay. Also, there was no consistency in the way this rule was applied. We saw plenty of other people with carry-ons larger than ours allowed on the plane and there was plenty of space in the overhead bins. 3. Every flight we were on was late. Two of three were late by over 45 mins. We didn't have any connections, but I imagine many connecting flights were missed. And, this didn't happen to us, but it's worth noting that when we landed in Athens, there were about 30 people waiting at the lost baggage desk. Most of them looked very angry. I asked one person what happened, and apparently an entire flight did not have their bags transferred at a connection. So not only will they force you to check your carry-on and then pay for a checked bag, they might lose it too. All-in-all, this is a pretty bad track record for 3 short island hop flights.	Couple Leisure	0	Economy Class
111	Aegean Airlines	Mytilini	Athens	NaN	Check-in uneventful. Boarding a zoo with a rude employee denying me boarding, claiming some zone story. Upon boarding I found a number of passengers that were probably the right 'zone' since they were seated in the rear already. An old ATR POS that Aegean must likely keep to connect some small islands that cannot accommodate jets. Why they would use these consistently for the third largest Greek island is unclear to me. Once seated in 15C and having put my backpack above, I see the flight attendant pushing it to make space for someone's big suitcase. When I objected, she pulled it down	Business	0	Economy Class

Airline Name	From	To	Via	Review	Type Of Traveller	Recommended	Seat Type
--------------	------	----	-----	--------	-------------------	-------------	-----------

and told me in the rudest way to put it under the seat in front of me. Said attendant did not wear a name tag.

```
In [15]: # Let's read one full review
print(df.loc[41].Review)
```

Excellent flights both ways with Adria Airways from Manchester to Ljubljana. Staff friendly but unobtrusive. Leg room far better than most economy flights, seats comfy and plenty of storage space. Flights on time, quick and efficient boarding and disembarking, smooth take offs and landings. All in all a very pleasant experience.

You try it

1. Use `str.contains` to find the number of reviews which mention the following keywords: 'Crew', 'Bag', 'Passenger'.
2. Calculate the total count across all 3 keywords.
3. Does review 186 from Aegean Airlines contain the word crew?

```
In [16]: for kw in ['Crew', 'Bag', 'Passenger']:
count = df.Review.str.contains(kw, case=False).sum()
print(f'{kw}: {count}')

total = sum(df.Review.str.contains(kw, case=False).sum() for kw in ['Crew', 'Bag', 'Passenger'])
print(f'Total: {total}')

# Does review 186 mention crew?
print(df.loc[186, 'Review'])
print('crew' in df.loc[186, 'Review'].lower())
```

Crew: 5202

Bag: 4174

Passenger: 4169

Total: 13545

They canceled my flight and did not notify me. I had to look the flight up myself to find out. Then they would only refund with a voucher that they won't send for 2 weeks and that expires in 12 months. So I will have already left Europe by the time I get it and will probably never be able to use it. I couldn't find a new flight within the same day because it was such short notice which screwed up plans for accommodation that I can't get refunded and so on. The stress and money they have cost me is crazy.


True

Topic Detection: Frequent Word Analysis

Advanced techniques like topic modeling (LDA) or transformer-based classification exist, but a simple and quick method is **word frequency analysis**. It gives a fast sense of recurring topics — service, food, delays, seats — without any machine learning.

```
In [17]: sample = df.sample(1000, random_state=42)
all_words = sample.Review.str.lower().str.split().sum()

pd.Series(all_words).value_counts().head(20)[::-1].plot.barh(color='orange');
```




The top terms are mostly **stopwords** — words that carry no meaning (the, a, I, was ...). We can filter them out using a standard list from NLTK:


```
In [18]: #!/ pip install nltk
import nltk
nltk.download('stopwords', quiet=True)
from nltk.corpus import stopwords

stop = set(stopwords.words('english'))

filtered_words = [w for w in all_words if w not in stop]
pd.Series(filtered_words).value_counts().head(20)[::-1].plot.barh(color='orange');
```




Let's try to plot this as a wordcloud. The beauty and power of python is that - regardless of your objective - you can always find some package that will get close to what you want:



```
In [19]: #!/ pip install wordcloud
from wordcloud import WordCloud
import matplotlib.pyplot as plt

text = ' '.join(filtered_words)
wordcloud = WordCloud(width=1200, height=400, background_color='#222222', colormap='summer').generate(text)


plt.figure(figsize=(15, 5))
plt.imshow(wordcloud, interpolation='bilinear')
plt.axis('off')
plt.show()
```

Here is an alternative way to write this using `.apply()`, which applies a function row-by-row. This pattern becomes essential for more complex scoring later:

```
In [22]: def sentiment(review):
is_good = 'good' in review.lower()
is_bad = 'bad' in review.lower()
return is_good - is_bad

df['simple_sentiment'] = df.Review.apply(sentiment)
df.simple_sentiment.hist(bins=5);
```



Litmus test

This all feels a bit ad-hoc and we should probably validate this with the recommendations to make sure the sentiment makes sense ... Does our simple sentiment score actually correlate with whether passengers recommend the airline?

```
In [23]: df[['simple_sentiment', 'Recommended']].corr()
```

```
Out[23]:
```

	simple_sentiment	Recommended
simple_sentiment	1.000000	0.319901
Recommended	0.319901	1.000000

Approach 2: Lexicon-Based (VADER)

The **VADER** (Valence Aware Dictionary and sEntiment Reasoner) lexicon was developed by Hutto & Gilbert (2014) at Georgia Tech. They had human annotators rate thousands of words for sentiment intensity, then built a dictionary capturing both polarity (positive/negative) and strength. It handles informal text, punctuation emphasis (!), and capitalization well.

```
In [24]: import nltk
from nltk.sentiment import SentimentIntensityAnalyzer

nltk.download('vader_lexicon', quiet=True)
sia = SentimentIntensityAnalyzer()
```

```
In [25]: # Peek at the lexicon – each word has a valence score
print(sia.lexicon.get("amazing"))
print(sia.lexicon.get("terrible"))
print(sia.lexicon.get("okay"))
```

```
2.8
-2.1
0.9
```

```
In [26]: # polarity_scores returns neg/neu/pos/compound – compound is the overall score (-1 to +1)
sia.polarity_scores('I really loved the snacks on the flight!')
```

```
Out[26]: {'neg': 0.0, 'neu': 0.572, 'pos': 0.428, 'compound': 0.6689}
```

```
In [27]: from tqdm import tqdm
tqdm.pandas()

def get_sentiment(review):
    return sia.polarity_scores(review)['compound']

df['sentiment'] = df['Review'].progress_apply(get_sentiment)
```

```
100%|██████████| 23046/23046 [00:14<00:00, 1557.23it/s]
```

You try it

1. Show the 2 most positive and 2 most negative reviews according to VADER.
2. Run the correlation test: does VADER sentiment correlate better with `Recommended` than the ad-hoc approach?

Hint: use `df.sort_values()`

```
In [28]: print('Most positive:')
print(df.sort_values('sentiment')[::-2].Review)

print('Most negative:')
print(df.sort_values('sentiment')[2:].Review)
```

Most positive:

13167 We heard about this new airline and decided to give it a try. The price seemed quite reasonable for what we were getting. This is an all business class airline. The price was not much more than flying economy on other airlines. We read the reviews and were aware of time changes and flight cancellations that have occurred in the past. This was not a business trip, so we were a bit flexible. We flew from Newark to London and then back a week later. Our flight out to London was delayed by about an hour, but we got plenty of notice and it did not affect it one bit. On the way home, our flight took off a couple of hours early, but again we got plenty of notice of the flight change. Truth is the early departure worked better for us, so we were not unhappy with the change. Check in went very smoothly. No lines and the staff were very courteous and friendly. We were traveling light, 2 suitcases between 3 of us, but it is nice to know that if we decided to buy something large, it could fly back with us as a piece of luggage. Once checked in we went straight for the lounge. Nothing special, but much better than sitting in the regular waiting area. You get some decent food and drink before you board. If you wait till the end, someone will come around and remind you when it is boarding time. As a business class passenger you get to take the short security line. Truth is that at both Luton and Newark, it did not make a big difference as lines at both airports were short. Apparently this airline has two Boeing 757 planes. We rode on both. On the way to London we rode on the newer plane. The seats are so much nicer than any I have been in before. They have a ton of legroom, I could not touch the seat in front of me. The seats go almost flat. The seat controls are not intuitive, but you can figure them out if you play with the controls enough. The most important controls are the ones that convert the seat to a bed and the one that converts it back to a seat. The newer 757 is nicer. It has more bathrooms and there are no wires to charge the entertainment tablets. The seats also end up a bit flatter in bed mode. Entertainment wise, you get a nice big Samsung tablet. On the older plane where the tablets need to be plugged in, quite a few were running out of battery during our flight. A bit annoying, but if you kept the tablet in the cradle in front of the seat and left it plugged in, it was not an issue. The selection of entertainment was a bit slim. There were 6 TV shows to watch and about a dozen movies. Lucky for me there were a couple of movies I had not yet seen so I was kept entertained. Food wise, no matter what this is airline food and it is not fine dining. With that said, they did the best they could. The fruit was very fresh and the deserts were really nice. Another nice touch was real silverware and cloth napkins. I'm not a drinker, so I can't comment on the wine and champagne offered on the flight. When I shop for anything, I look at value for the money spent, not just the cheapest price. Value wise I would rate this airline very high. There are only 76 passengers per flight, so you don't feel like cattle getting herded to the slaughter house. The staff is very well trained and friendly. If you don't mind the chance of a flight change, the value is hard to beat. I would use La Compagnie again.

16427

It's good to note that Qatar Airways open its check in counter 5 1/2 hours before departure! Good thing I went to NAIA cause I thought its going to rain and it might cause traffic! Started forming our lines at 5:30pm and counters open exactly 5:50pm for the 11:20pm flight. Check in process was quiet slow, despite my Silver Status I left the counter at 6:40pm. Understandably only two counters were open since the other check in staff are assisting the 6:35pm departure. Anyway, though it was slow but OK since the clerk was super nice and offered me a complimentary lounge pass! The plane departed around 15 minutes later than published but was able to overcome and arrived earlier in Doha! Was a very smooth landing and the crew were all warm! Food served dinner and breakfast were both sumptuous and delicious! I was seated on a window seat and was quiet cramped because the guy beside me was humungous. Arrival in Doha was as expected very warm and pleasant, security check was fast and quick! A separate lounge for Silver members were provided, and was surprised of its huge size, very comfortable, good food, very clean comfort rooms! Went to the gate early as it was quiet far but thanks to the sky train, I was at my gate in seconds! Fortunately, I was upgraded to Business Class on the second leg of the journey! It was a very pleasant experience considering the 13 hours flight to Montreal! Everything went perfectly well! The very comfortable Q Suite, the award winning Catering and the Best Crew Services contributed to an overall pleasant experience! Thanks to Anna and Dale for taking care of me plus their very beautiful Supervisor which I wasn't able to get the name! Thank you Qatar Airways for the free upgrade and thank you for the best service provided!

Name: Review, dtype: str

Most negative:

16936

I flew Business Class Bangkok-Amman; upon boarding, I could not find a space for my one piece of hand luggage, as the overhead locker and those near it were all full. There was a member of cabin crew looking at us, but he did not offer to help, so I asked him to help me find a space, as the economy class passengers were streaming through the business class cabin and I was finding it hard to move around. He seemed very unwilling, but I looked in a locker further away, so I mentioned that he could rearrange some backpacks on the locker above my seat or the one next to it to make some room. He told me not to tell him what to do and to shut up and wait, so I told him he was being rude, to which he responded by shouting at me in front of the entire Business Class cabin and moving menacingly towards me until he came face to face with me in a threatening manner, shouting at me. I was so taken aback and shocked, I told him not to raise his voice at me, that it was unnecessary and frankly wrong. He then called me the rude one and said he would have me taken off the flight if I didn't shut up. At this point my partner intervened and told him he had no right to talk to me in such a threatening manner, as I'd done nothing wrong, as in fact he was the one with the anger issue. The cabin crew member then shouted at my partner too, so I had to try and calm the situation down and asked the crew member to identify himself, as he wasn't wearing a name badge. He refused to answer, so I asked to speak to the senior cabin crew member, to which he replied that it was him! I therefore asked for his name and he refused to give it to me again. As space was found for my bag above my seat, I asked him to help me lift my bag into it: he refused and told me I have two hands and to do it myself! I pointed out that my visibly bandaged right hand was broken (I have a medical report to prove this too) and that it was an unkind and discriminatory thing to say to a passenger, to which he just laughed and stared at me with a sarcastic look. I then asked another crew member for help me, but he would not help me either saying he wasn't allowed to touch passenger's bags, even though I then saw him help several other passengers with their bags and jackets. Finally, my partner put my bag in the locker, and I sat down completely shocked and frightened by what had just happened. He told me that the crew member in question was indeed the senior member and that he should have been wearing a badge. You can imagine what a terrible flight it was: the senior crew member would strut down the aisle helping other passengers and bantering with them, whilst shooting me horrible, threatening looks: I was too scared to even eat my dinner, in case it'd been tampered with, as there was no trolley service whatsoever. We were only offered red wine, the food was awful and there was no trolley service and nobody asked if we wanted to drink or eat something more. The seats were old as well and the film selection very poor. When we disembarked, the senior member was at the door and refused to say goodbye; just awful. We are frequent travellers and have never come across such dreadful behaviour in our lives. We shall never fly with Royal Jordanian again.

2411 Cape Town to Windhoek. Never fly with Air Namibia, flew back to Frankfurt via Windhoek on 17.02.2019 after a four-week holiday in Cape Town. The problem actually started when checking in in Cape Town. We have never experienced such unfriendly and unwilling groundworkers who refuse work. The highlight of the outrage was, however, that for my golf baggage that the approved dimensions exceeded by as much as 6cm, a transport fee of 2 322, - Rand (about € 150, -) was charged. The checking-in employee put on his tape measure until he really came to the 6 cm. Of course we did not agree and tried by all means to avoid this arbitrary charge. Unfortunately, we have come to the wrong address with the two colleagues. Neither Ziupad nor Sean (Supervisor) were in any way willing to compromise. They shone through ignorance and unimaginable stubbornness, combined with a rudeness that is second to none. Even my objection that the luggage on the way in Frankfurt in any way appreciated and simply waved, the situation did not improve. After much discussion in the course of which I wanted to speak to the supervisor, they even did not shy away from letting the police come to the counter. We paid a total of € 4423.16 for the flight. And then this very nice holiday was completely messed up by this action. That was just the beginning of a flight full of mishaps. On the connecting flight to Frankfurt, I can only confirm what the travelers from France (N Ranelle) from 29.10.2018 had experienced before me. Very unfriendly crew. The two business class toilets were blocked for hours by the crew. Entertainment System failed in more than half of Business Class. No magazines. Seats only mediocre. Drink service during the night not available. The top, however, was the Business Lounge in Windhoek. Ripped sofas, makeshift louvered ceilings. Dirty carpeting, plastic cutlery and dishes. I have attached pictures. Unfortunately, I also have to say that the airline does not manage to keep any aircraft or building in order. Everything is used until it falls apart. In the future, we will fly to South Africa only Arab or Asian airlines, even if we will be traveling longer. For a stay of more than four weeks, this is not so significant and is more than offset by the friendly and perfect service, the modern airplanes and business lounges with smaller money. Leider auch bei uns: Nie mehr mit Air Namibia fliegen, wir sind am 17.02.2019 nach einem vier-wöchigen Urlaub in Kapstadt zurück nach Frankfurt über Windhoek geflogen. Das Problem ging eigentlich schon beim Einchecken in Kapstadt los. So ein unfreundliches und unwilliges, arbeitsverweigerndes Bodenpersonal haben wir noch nicht erlebt. Der Höhepunkt der Unverschämtheiten war aber der, dass mir für mein Golfgepäck, dass die zugelassenen Maße um sage und schreibe 6cm übertraf, eine Beförderungsgebühr von 2.322,- Rand (ca. € 150,-) in Rechnung gestellt wurde. Der eincheckende Mitarbeiter legte sein Maßband so oft an, bis er wirklich auf die 6 cm kam. Natürlich waren wir damit nicht einverstanden und versuchten mit allen Mitteln, diese willkürliche Gebühr zu vermeiden. Doch leider sind wir bei den beiden Kollegen an die falsche Adresse geraten. Weder Ziupad noch Sean (Supervisor) zeigten sich in irgendeiner Weise kompromissbereit. Sie glänzten durch Ignoranz und unvorstellbarer Sturheit, verbunden mit einer Unfreundlichkeit, die ihresgleichen sucht. Auch mein Einwand, dass das Gepäck auf dem Hinflug in Frankfurt in keiner Weise gewürdigt und einfach durchgewinkt wurde, verbesserte die Situation nicht. Nach heftiger Diskussion in deren Verlauf ich die Vorgesetzten sprechen wollte, scheuten sie sich sogar nicht, die Polizei zum Schalter kommen zu lassen. Wir haben für den Flug insgesamt € 4423,16 bezahlt. Und dann wurde uns dieser sehr schöne Urlaub durch diese Aktion komplett versaut. Das war nur der Anfang eines Fluges voller Pannen. Bei dem Anschlußflug nach Frankfurt kann ich nur bestätigen, was die Reisenden aus Frankreich (N Ranelle) vom 29.10.2018 vor mir erlebt hatten. Sehr unfreundliche Crew. Die beiden Business Class Toilettten wurden stundenlang von der Crew blockiert. Entertainment System bei mehr als der Hälfte der Business Class ausgefallen. Keine Zeitschriften. Sitze auch nur mittelmäßig. Getränkeservice während der Nacht nicht vorhanden. Die Spitze war allerdings die Business Lounge in Windhoek. Zerfetzte Sofas, notdürftig mit Decken kaschiert. Dreckiger Teppichboden, Plastik-Besteck und Geschirr. Ich habe Bilder beigefügt. Leider muß ich auch hier feststellen, dass es die Airline nicht schafft, weder Flugzeuge noch Gebäude in Stand zu halten. Es wird alles so lange benutzt, bis es auseinander fällt. Wir werden in Zukunft nach Südafrika nur noch arabischen oder asiatischen Airlines fliegen, auch wenn wir dadurch länger unterwegs sein werden. Bei einem Aufenthalt von mehr als vier Wochen fällt dies aber nicht so ins Gewicht und wird durch den freundlichen und perfekten Service, die modernen Flugzeuge und Business Lounges bei kleinerem Geld mehr als wett gemacht.

Name: Review, dtype: str

```
In [29]: print('\nCorrelation with Recommended:')
print(df[['sentiment', 'Recommended']].corr())
```

```
Correlation with Recommended:
          sentiment  Recommended
sentiment    1.000000    0.674943
Recommended    0.674943    1.000000
```

Approach 3: Machine Learning

A popular family of machine learning approaches to sentiment analysis is BERT (Bidirectional Encoder Representations from Transformers). Once you know how to use them, they become the swiss army knife of text analytics. BERT models have been fine-tuned for tasks like:

- Sentiment analysis
- Text classification
- Topic detection
- ...

For now, let's keep it simple with a lightweight version of BERT, already fine-tuned for sentiment: `distilbert-base-uncased-finetuned-sst-2-english`. It's fast, free to use, and good enough to show the power of transformer models over rule-based methods.

Note: this may not run on all computers, if it doesn't on yours, switch to google collab.

```
In [30]: #! pip install transformers torch
from transformers import pipeline
```

```
sentiment_model = pipeline(
    "sentiment-analysis",
```

```

model="distilbert-base-uncased-finetuned-sst-2-english",
truncation=True,
max_length=512
)

def scored_sentiment(text):
    result = sentiment_model(text)[0]
    return result['score'] if result['label'] == 'POSITIVE' else -result['score']

# Quick demo
print(scored_sentiment('Best. Flight. Ever.'))
print(scored_sentiment('Worst experience of my life. Never again.'))

```

```

Loading weights: 100%|██████████| 104/104 [00:00<00:00, 9217.86it/s]
0.9998503923416138
-0.9997391104698181

```

You try it

Use the `.progress_apply` function to apply sentiment scoring to all (or a sample of all) BERT sentiments.

```

In [31]: sample = df.sample(100, random_state=42)
bert_sentiments = sample.Review.progress_apply(scored_sentiment)

```

```

0%|          | 0/100 [00:00<?, ?it/s]
100%|██████████| 100/100 [00:10<00:00, 9.35it/s]

```

Going further: emotion detection

Binary positive/negative is useful but coarse. The model below - trained by your professor ;-) on the [GoEmotions dataset](#) - classifies text into 28 fine-grained emotions (joy, anger, disappointment, admiration ...). It's based on ModernBERT, a post-ChatGPT architecture, and available on [HuggingFace](#).

```

In [42]: from transformers import pipeline

classifier = pipeline(
    "text-classification",
    model="cirimus/modernbert-base-go-emotions",
    top_k=5
)

text = "The crew was stellar but my why mist my seat be so incredibly uncomfortable for a 12-hour flight."
predictions = classifier(text)

top_5 = sorted(predictions[0], key = lambda x: x["score"], reverse = True)[:5]
print("Top 5 emotions detected:")
for pred in top_5:
    print(f" {pred['label']:15s}: {pred['score']:.3f}")

```

```


Loading weights: 100%|██████████| 138/138 [00:00<00:00, 5907.11it/s]
Top 5 emotions detected:
admiration      : 0.641
disappointment  : 0.271
annoyance       : 0.096
neutral         : 0.079
disapproval     : 0.055

```

```

In [50]: df.groupby('Airline Name')['sentiment'].mean().sort_values().plot.barh(
    color='orange', figsize=(8,12)
);

```




Who Are the Stars in Star Alliance?

Now that we have sentiment scores, let's rank airlines. We'll filter to Star Alliance members for a focused comparison.

```
In [45]: star_alliance_members = [
    'Aegean Airlines', 'Air Canada', 'Air China', 'Air India',
    'ANA All Nippon Airways', 'Asiana Airlines', 'Austrian Airlines',
    'Avianca', 'Brussels Airlines', 'Copa Airlines', 'Croatia Airlines',
    'Egyptair', 'Ethiopian Airlines', 'Eva Air', 'LOT Polish Airlines',
    'Lufthansa', 'Scandinavian Airlines', 'Shenzhen Airlines',
    'Singapore Airlines', 'South African Airways',
    'Swiss International Air Lines', 'TAP Portugal',
    'Thai Airways', 'Turkish Airlines', 'United Airlines'
]

f_star = df['Airline Name'].isin(star_alliance_members)
df[f_star].groupby('Airline Name')['sentiment'].mean().sort_values().plot.barh(
    color='steelblue', title='Star Alliance - Average VADER Sentiment'
);
```

What follow-up analysis would you do if you worked at Star Alliance?

Part III: Aspect-Based Analysis with LLMs

Overall sentiment tells you whether a review is positive or negative. But an airline category manager needs sharper answers: *How is the crew rated? What do passengers say about luggage? How does food score on long-haul routes?*

This is **aspect-based sentiment analysis** — and it's where LLMs shine. Rather than training a separate model for each aspect (crew, food, luggage, punctuality ...), you just change the prompt.

You have two options for any custom NLP task:

1. Fine-tune your own model (requires labelled data — take the ML class)
2. **Use an LLM to annotate your data** — powerful, fast, and requires no training data

Let's try the second with Google Gemini.

Setup:

1. Create a [Google account](#) or log in.
2. Create a free [Gemini API key](#).
3. Install the SDK: `pip install -q -U google-generativeai`

```
In [47]: API_KEY = '...' # paste your key here
```

```
In [50]: # ! pip install -q -U google-generativeai
from google import generativeai

client = generativeai.Client(api_key=API_KEY)

response = client.models.generate_content(
    model="gemini-2.5-flash",
    contents="Explain in one number what the meaning of life is."
)
print(response.text)
```

42

Scoring an aspect: crew

We'll start by scoring the **crew** aspect — one of the most frequently mentioned in airline reviews. The key idea: write a prompt that extracts a structured score from free text. First test your prompt manually in your favourite LLM chat interface, then automate it below.

```
In [51]: df_ana = df[df['Airline Name'] == 'ANA All Nippon Airways']
df_ana[['Airline Name', 'Review']].head()
```


Out[51]:	Airline Name	Review
4350	ANA All Nippon Airways	Pretty disappointed with ANA, especially when compared to JAL. Seats are dated and half are very exposed. The Japanese food was generally good and at times excellent. The champagne was undrinkable. The pre-flight drinks (where offered) were from plastic disposable cups. The main listed wine selection on every leg was unavailable. The crew were generally good and attentive. It is such a pity that they appear to be economising.
4351	ANA All Nippon Airways	Great service, lovely crew, comfortable seats, lots of leg room (I am 6'2"), roomy toilets and an immaculately clean aircraft. It's hard to imagine how a 14 hour flight in economy could be done any better. I paid in advance an extra £15 for a special Japanese meal which was most definitely worth it and really delicious, beautifully presented.
4352	ANA All Nippon Airways	Overall was a good experience, check in went smoothly, food was decent for airline food and the staff was friendly. The main downside was that the lounge in Tokyo Haneda was overcrowded.
4353	ANA All Nippon Airways	Paid for seats so entire family is together. Airline changes aircraft/seats and puts a stranger between us. Not even a single thought goes in this process. Extremely difficult to connect with someone on the phone. Unfortunately the competition is worse. Over the years flying changed from a fun way to travel to a dreadful experience.
4354	ANA All Nippon Airways	I'm super unhappy about ANA service. They asked me and my daughter to move our luggage after we have settled down, me and my daughter are both short, but she didn't offer her help, so that I almost dropped my luggage on one of passengers. Later during the breakfast time, she didn't offer food to my daughter until I figured it out. My daughter is shy girl, so I had to step out and ask her to bring the breakfast to her. This kind of things have never happened to me before. I'm extremely unhappy about ANA service.

In [52]: `PROMPT = ""`
 You're a world-class customer service analyst who is analyzing airline reviews.
 Score the crew on the flight (friendliness, professionalism, helpfulness) from 0 to 10.
 If crew is not mentioned, return 'N/A'.

Also extract the most relevant passage that supports your score.

Format your response as:

`<score>YOUR_SCORE_HERE</score><passage>YOUR_PASSAGE_HERE</passage>`

Important: only reply with the required tags and scores. Don't add anything else.
 ""

```
def score_crew(review_text):
    prompt = PROMPT + "\nReview:" + review_text
    response = client.models.generate_content(
        model="gemini-2.5-flash",
        contents=prompt
    )
    return response.text.strip()

score_crew('the crew was amazing. the food was good too.')
```

Out[52]: `'<score>10</score><passage>the crew was amazing.</passage>'`

You try it

1. Apply `score_crew()` to `df_sample` below (25 rows - don't run on the full dataset yet, it uses API quota).
2. Parse the output: create a `CrewScore` column and a `Rationale` column.
3. **Bonus:** change the `PROMPT` to score a *different* aspect - luggage handling, food quality, or punctuality. How little do you need to change to pivot to a completely different analysis?

Hints:

- Use `.apply(score_crew)` or `.progress_apply(score_crew)` to get a column of raw LLM output strings.
- Use string splitting on `<score>` and `<passage>` tags to extract the values.
- LLMs sometimes don't follow instructions ... filter out rows where the tags are missing.

In [53]: `df_sample = df.sample(25, random_state=42)`
`# llm_scores = df_sample['Review'].apply(score_crew)`
`# ...`